

MSBD566_Predictive_Final

December 12, 2025

0.1 MSBD566 - Final Project

0.2 Esmeralda Garcia

Data Source: https://www.kaggle.com/datasets/samiraalipour/genomics-of-drug-sensitivity-in-cancer-gdsc?resource=download&select=GDSC_DATASET.csv

```
[8]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.decomposition import PCA
from sklearn.metrics import classification_report, confusion_matrix

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.utils import to_categorical

pd.set_option("display.max_columns", 100)
plt.rcParams["figure.dpi"] = 120

print("TensorFlow version:", tf.__version__)
```

TensorFlow version: 2.11.0

```
[9]: df = pd.read_csv("GDSC_DATASET.csv")
print("Raw shape:", df.shape)
df.head()

LOW, HIGH = 0.0, np.log(10.0)

def label_ic50(x):
    if x < LOW:
        return "Very potent"
```

```

    elif x <= HIGH:
        return "Average sensitivity"
    else:
        return "Resistant"

df = df.dropna(subset=["LN_IC50"]).copy()
df["TargetClass"] = df["LN_IC50"].apply(label_ic50)

df["TargetClass"].value_counts()

```


Raw shape: (242035, 19)

[9]: TargetClass

Resistant	157188
Average sensitivity	51171
Very potent	33676

Name: count, dtype: int64

[10]: ax = df["TargetClass"].value_counts().reindex([
 "Very potent", "Average sensitivity", "Resistant"])
 .plot(kind="bar")
 ax.set_title("Distribution of Drug Sensitivity Classes")
 ax.set_xlabel("Class")
 ax.set_ylabel("Count")
 plt.xticks(rotation=15)
 plt.tight_layout()
 plt.show()



```
[5]: X_num.dtypes
```

```
X_num.apply(lambda col: pd.to_numeric(col, errors="coerce")).isna().sum()
```

```
[5]: AUC          0
Z_SCORE       0
CNA         242035
Gene Expression 242035
Methylation   242035
dtype: int64
```

```
[11]: num_feats = [c for c in ["AUC", "Z_SCORE"] if c in df.columns]
```

```
cat_feats = [
    "GDSC Tissue descriptor 1",
    "TARGET_PATHWAY",
    "Microsatellite instability Status (MSI)",
    "Screen Medium",
    "Growth Properties",
    "CNA",
    "Gene Expression",
```

```

    "Methylation"
]

cat_feats = [c for c in cat_feats if c in df.columns]

print("Numeric features:", num_feats)
print("Categorical features:", cat_feats)

```

Numeric features: ['AUC', 'Z_SCORE']
Categorical features: ['GDSC Tissue descriptor 1', 'TARGET_PATHWAY',
'Microsatellite instability Status (MSI)', 'Screen Medium', 'Growth Properties',
'CNA', 'Gene Expression', 'Methylation']

```
[12]: X_num = df[num_feats].copy()
X_cat = df[cat_feats].copy()
y = df["TargetClass"].copy()

print("X_num shape:", X_num.shape)
print("X_cat shape:", X_cat.shape)
print("y shape:", y.shape)
```

X_num shape: (242035, 2)
X_cat shape: (242035, 8)
y shape: (242035,)

```
[15]: X_num_clean = X_num.apply(lambda col: pd.to_numeric(col, errors="coerce"))

num_imputer = SimpleImputer(strategy="median")
X_num_imputed = num_imputer.fit_transform(X_num_clean)

scaler = StandardScaler()
X_num_scaled = scaler.fit_transform(X_num_imputed)

print("Numeric scaled shape:", X_num_scaled.shape)

cat_imputer = SimpleImputer(strategy="most_frequent")
X_cat_imputed = cat_imputer.fit_transform(X_cat)

ohe = OneHotEncoder(handle_unknown="ignore", sparse_output=False)
X_cat_encoded = ohe.fit_transform(X_cat_imputed)


print("Categorical encoded shape:", X_cat_encoded.shape)
```

Numeric scaled shape: (242035, 2)
Categorical encoded shape: (242035, 56)

```
[16]: pca_full = PCA()
pca_full.fit(X_num_scaled)

explained_var_ratio = pca_full.explained_variance_ratio_
cum_explained = np.cumsum(explained_var_ratio)

plt.figure()
plt.plot(range(1, len(cum_explained)+1), cum_explained, marker="o")
plt.axhline(0.9, color="red", linestyle="--", label="90% variance")
plt.xlabel("Number of Principal Components")
plt.ylabel("Cumulative Explained Variance")
plt.title("PCA: Cumulative Explained Variance (Numeric Features)")
plt.legend()
plt.tight_layout()
plt.show()
```



```
[17]: n_components = 2

pca = PCA(n_components=n_components)
X_num_pca = pca.fit_transform(X_num_scaled)
```

```

print("Original numeric shape:", X_num_scaled.shape)
print("After PCA numeric shape:", X_num_pca.shape)
print("Explained variance of chosen components:", pca.explained_variance_ratio_)
print("Total variance explained:", pca.explained_variance_ratio_.sum())

```

Original numeric shape: (242035, 2)
After PCA numeric shape: (242035, 2)
Explained variance of chosen components: [0.72129403 0.27870597]
Total variance explained: 1.0

[18]: X_all = np.hstack([X_num_pca, X_cat_encoded])
print("Final feature matrix shape:", X_all.shape)

Final feature matrix shape: (242035, 58)

[19]: X_train, X_test, y_train_text, y_test_text = train_test_split(
X_all, y, test_size=0.2, random_state=42, stratify=y
)

print("Train size:", X_train.shape[0])
print("Test size:", X_test.shape[0])
print("Train class counts:\n", y_train_text.value_counts())
print("\nTest class counts:\n", y_test_text.value_counts())

label_to_int = {"Very potent": 0, "Average sensitivity": 1, "Resistant": 2}
int_to_label = {v: k for k, v in label_to_int.items()}

y_train_int = y_train_text.map(label_to_int).values
y_test_int = y_test_text.map(label_to_int).values

y_train_oh = to_categorical(y_train_int, num_classes=3)
y_test_oh = to_categorical(y_test_int, num_classes=3)

print("y_train_int shape:", y_train_int.shape)
print("y_train_oh shape:", y_train_oh.shape)

Train size: 193628
Test size: 48407
Train class counts:
TargetClass
Resistant 125750
Average sensitivity 40937
Very potent 26941
Name: count, dtype: int64

Test class counts:
TargetClass
Resistant 31438

```
Average sensitivity      10234
Very potent              6735
Name: count, dtype: int64
y_train_int shape: (193628,)
y_train_oh shape: (193628, 3)
```

```
[20]: input_dim = X_train.shape[1]
```

```
model = Sequential([
    Dense(64, activation="relu", input_shape=(input_dim,)),
    Dropout(0.3),
    Dense(32, activation="relu"),
    Dropout(0.2),
    Dense(3, activation="softmax")
])

model.compile(
    loss="categorical_crossentropy",
    optimizer="adam",
    metrics=["accuracy"]
)

model.summary()
```

```
history = model.fit(
    X_train, y_train_oh,
    validation_split=0.2,
    epochs=20,
    batch_size=256,
    verbose=1
)
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
<hr/>		
dense (Dense)	(None, 64)	3776
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 32)	2080
dropout_1 (Dropout)	(None, 32)	0
dense_2 (Dense)	(None, 3)	99
<hr/>		
Total params: 5,955		

```
Trainable params: 5,955
Non-trainable params: 0
```

```
2025-12-07 23:59:43.533800: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations: SSE4.1 SSE4.2 AVX AVX2 AVX512F FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
2025-12-07 23:59:43.538504: I tensorflow/core/common_runtime/process_util.cc:146] Creating new thread pool
with default inter op setting: 2. Tune using inter_op_parallelism_threads for
best performance.
```

```
Epoch 1/20
606/606 [=====] - 7s 10ms/step - loss: 0.4802 -
accuracy: 0.8143 - val_loss: 0.3875 - val_accuracy: 0.8466
Epoch 2/20
606/606 [=====] - 4s 6ms/step - loss: 0.4023 -
accuracy: 0.8432 - val_loss: 0.3728 - val_accuracy: 0.8499
Epoch 3/20
606/606 [=====] - 5s 8ms/step - loss: 0.3859 -
accuracy: 0.8478 - val_loss: 0.3628 - val_accuracy: 0.8505
Epoch 4/20
606/606 [=====] - 5s 8ms/step - loss: 0.3771 -
accuracy: 0.8493 - val_loss: 0.3578 - val_accuracy: 0.8522
Epoch 5/20
606/606 [=====] - 13s 21ms/step - loss: 0.3707 -
accuracy: 0.8510 - val_loss: 0.3532 - val_accuracy: 0.8541
Epoch 6/20
606/606 [=====] - 6s 10ms/step - loss: 0.3644 -
accuracy: 0.8531 - val_loss: 0.3475 - val_accuracy: 0.8568
Epoch 7/20
606/606 [=====] - 5s 9ms/step - loss: 0.3609 -
accuracy: 0.8537 - val_loss: 0.3434 - val_accuracy: 0.8561
Epoch 8/20
606/606 [=====] - 6s 9ms/step - loss: 0.3574 -
accuracy: 0.8547 - val_loss: 0.3418 - val_accuracy: 0.8579
Epoch 9/20
606/606 [=====] - 6s 9ms/step - loss: 0.3546 -
accuracy: 0.8559 - val_loss: 0.3389 - val_accuracy: 0.8595
Epoch 10/20
606/606 [=====] - 6s 10ms/step - loss: 0.3519 -
accuracy: 0.8567 - val_loss: 0.3358 - val_accuracy: 0.8597
Epoch 11/20
606/606 [=====] - 5s 7ms/step - loss: 0.3491 -
accuracy: 0.8575 - val_loss: 0.3344 - val_accuracy: 0.8600
Epoch 12/20
```


```


606/606 [=====] - 4s 6ms/step - loss: 0.3483 -
accuracy: 0.8568 - val_loss: 0.3323 - val_accuracy: 0.8610
Epoch 13/20
606/606 [=====] - 4s 6ms/step - loss: 0.3464 -
accuracy: 0.8589 - val_loss: 0.3323 - val_accuracy: 0.8614
Epoch 14/20
606/606 [=====] - 4s 6ms/step - loss: 0.3440 -
accuracy: 0.8588 - val_loss: 0.3295 - val_accuracy: 0.8629
Epoch 15/20
606/606 [=====] - 5s 9ms/step - loss: 0.3431 -
accuracy: 0.8588 - val_loss: 0.3283 - val_accuracy: 0.8633
Epoch 16/20
606/606 [=====] - 9s 14ms/step - loss: 0.3429 -
accuracy: 0.8593 - val_loss: 0.3269 - val_accuracy: 0.8624
Epoch 17/20
606/606 [=====] - 4s 6ms/step - loss: 0.3409 -
accuracy: 0.8596 - val_loss: 0.3267 - val_accuracy: 0.8633
Epoch 18/20
606/606 [=====] - 4s 6ms/step - loss: 0.3404 -
accuracy: 0.8604 - val_loss: 0.3294 - val_accuracy: 0.8628
Epoch 19/20
606/606 [=====] - 4s 7ms/step - loss: 0.3393 -
accuracy: 0.8600 - val_loss: 0.3258 - val_accuracy: 0.8642
Epoch 20/20
606/606 [=====] - 4s 6ms/step - loss: 0.3391 -
accuracy: 0.8608 - val_loss: 0.3241 - val_accuracy: 0.8640

```

```
[21]: plt.figure()
plt.plot(history.history["accuracy"], label="Train accuracy")
plt.plot(history.history["val_accuracy"], label="Val accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title("Training vs Validation Accuracy (Neural Network)")
plt.legend()
plt.tight_layout()
plt.show()

plt.figure()
plt.plot(history.history["loss"], label="Train loss")
plt.plot(history.history["val_loss"], label="Val loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Training vs Validation Loss (Neural Network)")
plt.legend()
plt.tight_layout()
plt.show()
```





```
[22]: test_loss, test_acc = model.evaluate(X_test, y_test_oh, verbose=0)
print(f"Test accuracy: {test_acc:.3f}")
print(f"Test loss: {test_loss:.3f}")

y_pred_proba = model.predict(X_test)
y_pred_int = np.argmax(y_pred_proba, axis=1)
y_pred_labels = pd.Series(y_pred_int).map(int_to_label)

print("Classification report (Neural Network):")
print(classification_report(y_test_text, y_pred_labels, digits=3))
```

Test accuracy: 0.864
Test loss: 0.323
1513/1513 [=====] - 10s 7ms/step
Classification report (Neural Network):


	precision	recall	f1-score	support
Average sensitivity	0.744	0.672	0.706	10234
Resistant	0.893	0.961	0.926	31438
Very potent	0.891	0.705	0.788	6735

accuracy			0.864	48407
macro avg	0.843	0.779	0.806	48407
weighted avg	0.861	0.864	0.860	48407

```
[23]: labels_order = ["Very potent", "Average sensitivity", "Resistant"]

cm = confusion_matrix(y_test_text, y_pred_labels, labels=labels_order)
cm_df = pd.DataFrame(cm, index=labels_order, columns=labels_order)
cm_perc = cm_df.div(cm_df.sum(axis=1), axis=0) * 100

plt.figure(figsize=(7,6))
sns.heatmap(cm_perc, annot=True, fmt=".1f", cmap="Blues",
            xticklabels=labels_order, yticklabels=labels_order,
            cbar_kws={'label': '% of True Class'})
plt.title("Neural Network Confusion Matrix (% by True Class)")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.tight_layout()
plt.show()
```



[]: